

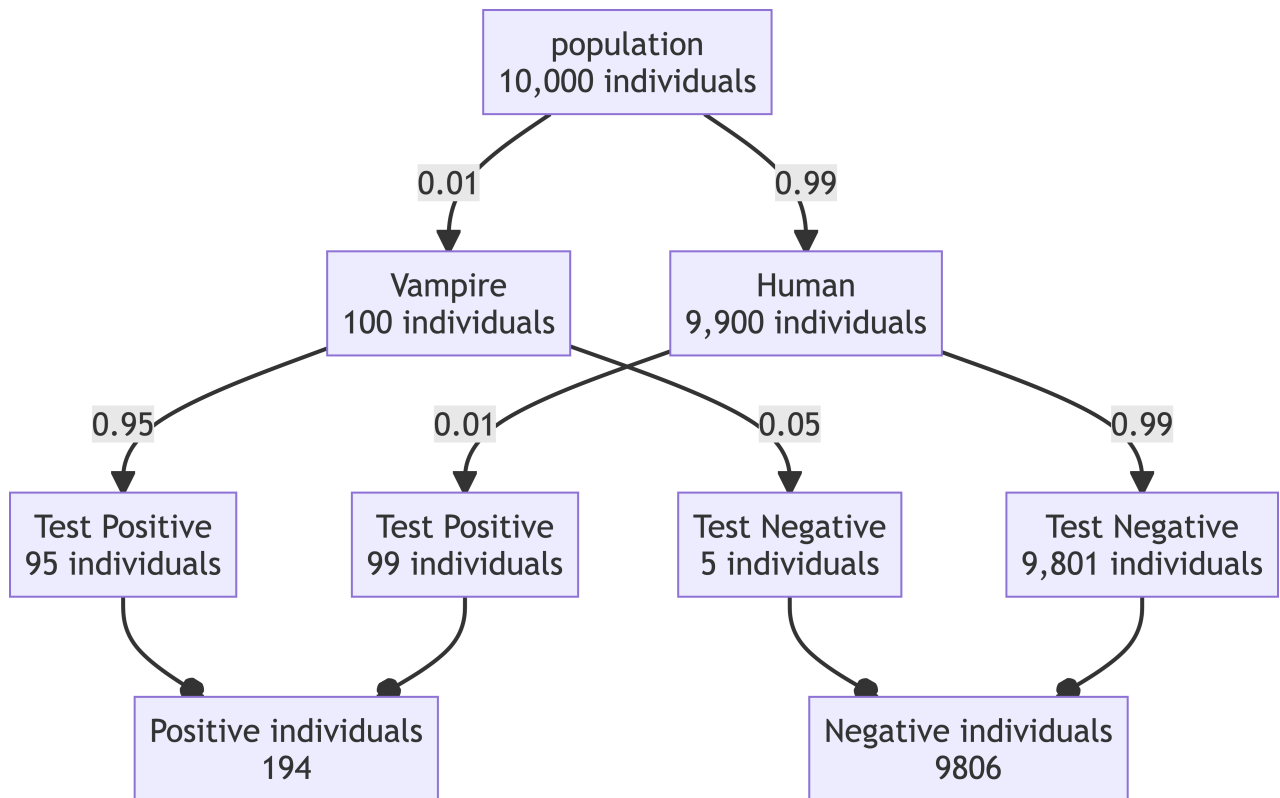
Starting Sampling

Josef Fruehwald

 Listening

```
library(tidyverse)
library(patchwork)
library(here)
library(gt)
source(here("_defaults.R"))
```

Classic Base Rate Issues



Plot of the base rate vs $P(\text{vampire} \mid \text{positive test})$

```
tibble(  
  # might as well get logarithmic  
  base_rate = 10^(seq(-3, -1, length = 20)),  
  vamp_and_pos = base_rate * 0.95,  
  vamp_and_neg = base_rate * 0.05,  
  human_and_pos = (1-base_rate) * 0.01,  
  human_and_neg = (1-base_rate) * 0.99,  
  p_vamp_pos = vamp_and_pos/(vamp_and_pos + human_and_pos),  
  p_hum_neg = human_and_neg/(vamp_and_neg + human_and_neg)  
) -> test_metrics
```

```
test_metrics |>  
  ggplot(aes(base_rate, p_vamp_pos))+
```

```

geom_point(color = "steelblue",
           size = 3)+
geom_line(color = "steelblue",
          linewidth = 1)+
scale_x_log10()+
ylim(0,1)+
labs(x = "P(vampire)",
     y = "P(vampire | positive)",
     subtitle = "P(positive | vampire) = 0.95\nP(positive | human) = 0.01",
     title = "Positive Predictive Value") +
theme(plot.subtitle = element_text(size = 12))

```




Figure 1: Probability someone *is* a vampire, given that they tested positive, relative to the base rate of being a vampire

```

test_metrics |>
ggplot(aes(base_rate, p_hum_neg))+
  geom_point(color = "steelblue",
            size = 3)+
  geom_line(color = "steelblue",
            linewidth = 1)+

```

```

scale_x_log10()+
labs(x = "P(vampire)",
     y = "P(human | negative)",
     subtitle = "P(positive | vampire) = 0.95\nP(positive | human) = 0.01",
     title = "Negative Predictive Value") +
theme(plot.subtitle = element_text(size = 12))

```




Figure 2: Probability of being a human given a negative test, relative to the base rate of being a vampire.

Tibble grid sampling

```

grid <- tibble(
  prob = seq(0.0001, 0.9999, length = 5000),
  prior_unstd = case_when(
    prob < 0.5 ~ 0,
    .default = 1
  ),
  prior_std = prior_unstd/sum(prior_unstd),
  data = dbinom(6, size = 9, prob = prob),

```

```

posterior_unstd = prior_std * data,
posterior = posterior_unstd / sum(posterior_unstd)
)

grid |>
  ggplot(aes(prob, prior_std))+
    geom_line()+
    labs(y = "prior density",
         title = "Prior") ->
  prior_plot

grid |>
  ggplot(aes(prob, data))+
    geom_line()+
    labs(y = "data density",
         title = "Data") ->
  data_plot

grid |>
  ggplot(aes(prob, posterior))+
    geom_line() +
    labs(y = "posterior density",
         title = "Posterior") ->
  posterior_plot

prior_plot | data_plot | posterior_plot

```




Figure 3: Prior, Data, Posterior


Sampling from the posterior

```
grid |>
  sample_n(size = 1e4,
            replace = T,
            weight = posterior)->
  posterior_samples

renv::install("tidybayes")

library(tidybayes)

posterior_samples |>
  ggplot(aes(prob))+
    stat_halfeye(point_interval = median_hdc)+
    geom_line(aes(y = posterior/max(posterior)),
              linetype = 2)+
    xlim(0,1)+
    theme(
      axis.title.y = element_blank(),
      axis.text.y = element_blank(),
      panel.grid.major.y = element_blank()
    )
```



```
posterior_samples |>
  median_hdci(prob, .width = c(0.5, 0.95)) |>
  gt() |>
  fmt_number(decimals = 2)
```

prob	.lower	.upper	.width	.point	.interval
0.67	0.58	0.74	0.50	median	hdci
0.67	0.50	0.86	0.95	median	hdci

BRMS

```
library(brms)

tibble(
  water = 6,
  samples = 9
)->
  water_to_model
```

```

water_form <- bf(
  water | trials(samples) ~ 1,
  family = binomial(link = "identity")
)

brm(
  water | trials(samples) ~ 1,
  data = water_to_model,
  family = binomial(link = "identity"),
  prior(beta(1, 1), class = Intercept, ub = 1, lb = 0),
  file_refit = "on_change",
  file = "water_fit.rds"
) ->
  water_model

```

```
water_model
```

```

Family: binomial
Links: mu = identity
Formula: water | trials(samples) ~ 1
Data: water_to_model (Number of observations: 1)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000

```

Population-Level Effects:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	0.64	0.14	0.35	0.88	1.00	1598	1945

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```

library(gtsummary)

water_model |>
  gtsummary::tbl_regression(intercept = T)

```


Characteristic	Beta	95% CI
(Intercept)	0.64	0.35, 0.88

```
fs::dir_create("tables")
```

```
water_model |>
  spread_draws(b_Intercept) |>
  summarise(Intercept = list(b_Intercept)) |>
  gt() |>
  gtExtras::gt_plt_dist(Intercept) -> modtable

modtable |>
  gtsave("tables/model.png")
```